

Filtering Microstructural Toxicity: A Spence-Harsanyi Signaling Approach to Regime-Routed Portfolios

Vishu Chauhan

University of Delhi

Department of Business Economics, South Campus

Abstract

For decades, algorithmic portfolio allocation has relied on continuous factor models and linear statistical arbitrage. However, post-2020 macroeconomic shocks have demonstrated that traditional strategies frequently succumb to retail-driven order-flow toxicity and false breakouts. In this paper, we propose a novel "Unified Macro-Micro" quantitative architecture. We utilize an unsupervised Hidden Markov Model (HMM) as a macro-regime router to dynamically allocate capital across a multi-strategy portfolio. Crucially, we introduce a microstructural veto layer grounded in Spence's Costly Signaling Theory and Harsanyi Bayesian updates. We formalize a Relative Volume-Weighted Spread (RVWS) metric—benchmarked against an institutional proxy of high volume and tight bid-ask spreads—to mathematically calculate the execution cost of price action. Through a 5-year rolling out-of-sample backtest incorporating a strict 5 basis point transaction cost across a cross-sectional sector basket, we demonstrate that a static Bayesian constraint sacrifices continuous trade frequency to generate a highly convex return profile. The strict model yielded a 40.27% net cumulative return while mitigating maximum drawdown to 1.86% during historic market stress. Furthermore, benchmarking indicates that RVWS demonstrates superior performance within the tested sample compared to standard volumetric imbalance metrics (VPIN) by successfully filtering toxic meme-stock anomalies. The architecture generated a Calmar Ratio of 21.65 (vs. 2.90 for the S&P 500) and a Sortino Ratio of 5.11, providing strong empirical evidence for its utility in capital-constrained institutional mandates requiring tail-risk immunization.

1 Introduction

The fundamental challenge in modern quantitative finance is no longer signal generation, but rather regime awareness and order-flow validation. The prolifera-

tion of zero-interest-rate monetary policies followed by historic inflation shocks has permanently altered the structural dynamics of asset pricing.

The integration of multi-strategy architectures often fails because macroscopic state models and microstructural execution models operate in isolated academic silos. A portfolio manager may correctly identify a trending macroeconomic regime using state-space models, but execute a trade based on a false breakout driven entirely by uninformed noise traders.

This paper bridges the gap between macroscopic state detection and microstructural game theory. We hypothesize that by treating market microstructure as a game of incomplete information (Harsanyi, 1967), we can mathematically quantify the financial "cost" of a price movement. By requiring a prohibitively high signaling cost to validate institutional presence, we construct a universal Bayesian veto layer designed to mitigate a portfolio's exposure to retail toxicity.

2 Literature Review

2.1 Regime-Switching Portfolios

The application of Hidden Markov Models (HMMs) to financial time series was pioneered by Hamilton (1989) and extended into portfolio allocation by Ang and Bekaert (2002). While HMMs are standard infrastructure for identifying macro-states, recent literature demonstrates they are insufficient as standalone execution triggers due to high false-positive rates in volatile microstructures (Kim et al., 2019).

2.2 Microstructural Toxicity

Existing microstructure literature heavily relies on probabilistic volume metrics to gauge informed trading. Easley, López de Prado, and O'Hara (2011) introduced Volume-Synchronized Probability of Informed Trading (VPIN) as an indicator of order-flow toxicity. While VPIN measures volumetric imbalances, it does not inherently calculate the capital expenditure

required to generate that imbalance. Our approach complements Kyle’s (1985) models of continuous auction insider trading and Glosten and Milgrom’s (1985) adverse selection models by framing order flow toxicity as a separating equilibrium problem.

2.3 Signaling Theory in Finance

Spence (1973) established that in markets with information asymmetry, informed agents can separate themselves from uninformed agents by taking actions that are prohibitively costly for the uninformed to mimic. Its direct algorithmic application to daily order flow as a binary execution veto remains under-explored.

3 Theoretical Framework

3.1 The Macro Router (HMM)

Let the global market state at time t be represented by an unobserved discrete variable $S_t \in \{1, 2, 3\}$, corresponding to Calm, Trending, and Crisis regimes. The HMM was trained using daily log returns and realized variance utilizing a 252-day rolling window to eliminate look-ahead bias. The Bayesian Information Criterion (BIC) selected a 3-state model.

Table 1: HMM Estimated Transition Probability Matrix

Current State (t)	Calm ($t + 1$)	Trend ($t + 1$)	Crisis ($t + 1$)
Calm	0.942	0.041	0.017
Trending	0.058	0.884	0.058
Crisis	0.012	0.063	0.925

3.2 The Micro Veto Layer and Economic Intuition

In market microstructure, price does not move simply because there are “more buyers than sellers”; it moves when aggressive market orders consume passive limit orders, sweeping the order book. Retail traders generally trade at the margin and cannot absorb deep liquidity. Therefore, a price movement requiring massive capital expenditure is the literal footprint of institutional urgency.

To quantify this, we define a market participant as type $\theta \in \{I, U\}$, where I is an Informed institution and U is an Uninformed retail trader. We define the observable signal s_t as the Relative Volume-Weighted Spread (RVWS), normalized over a 20-day rolling

window:

$$s_t = \frac{|C_t - O_t| \times V_t}{\frac{1}{20} \sum_{i=1}^{20} (|C_{t-i} - O_{t-i}| \times V_{t-i})} \quad (1)$$

To achieve a separating equilibrium, the cost function $c(s, \theta)$ must satisfy the single-crossing property:

$$\frac{\partial c(s, U)}{\partial s} > \frac{\partial c(s, I)}{\partial s} \quad (2)$$

To empirically verify this property, we tracked the $T + 5$ profitability across RVWS cohorts (Table 2).

Table 2: Single-Crossing Property ($T + 5$ Win Rate)

RVWS Range	Inst. Proxy	Retail Proxy
0.5 – 1.0	51.2%	49.8%
1.0 – 2.0	54.7%	46.1%
2.0 – 3.0	59.3%	41.2%
> 3.0	63.8%	35.4%

3.3 Empirical Harsanyi Bayesian Updater

To operationalize the Bayesian update, we define our institutional proxy (I) as days where $V_t > \mu_V + 1.5\sigma_V$ **and** the intraday bid-ask spread is below the 50th percentile. Retail flow (U) is proxied by high volume paired with top-quartile spread widening.

The uninformative prior $P(I)$ was set to 0.20. A Jarque-Bera test strongly rejected normality ($p < 0.01$). Consequently, we model $P(s_t|\theta)$ using a Student- t distribution with ν degrees of freedom:

$$P(s_t|\theta) = \frac{\Gamma(\frac{\nu+1}{2})}{\Gamma(\frac{\nu}{2})\sqrt{\pi\nu}\sigma_\theta} \left(1 + \frac{(s_t - \mu_\theta)^2}{\nu\sigma_\theta^2}\right)^{-\frac{\nu+1}{2}} \quad (3)$$

The algorithmic engine updates its prior belief using Bayes’ rule:

$$P(I|s_t) = \frac{P(s_t|I)P(I)}{P(s_t|I)P(I) + P(s_t|U)P(U)} \quad (4)$$

Execution is constrained by a strict verification threshold, τ :

$$\text{Execution}_t = \begin{cases} 1, & \text{if } P(I|s_t) > \tau \\ 0, & \text{otherwise} \end{cases} \quad (5)$$

4 Experimental Design

An out-of-sample backtest was conducted (2021–2026). The asset universe included SPY, GLD/SLV, VIXY,

Table 3: 5-Year Out-of-Sample Sector Basket Metrics (Net of 5 bps TC)

Metric	Model A (Strict $\tau=0.6$)	Model B (Adaptive)	Model D (VPIN)	SPY (Benchmark)
Cumulative Return	40.27%	-28.95%	18.40%	73.81%
Max Drawdown	1.86%	38.57%	14.20%	25.40%
Sortino Ratio	5.11	-0.58	0.95	0.88
Return Skewness	+35.34	-1.12	+1.40	-0.68
Calmar Ratio	21.65	-0.75	1.29	2.90
Trade Count	12	642	145	N/A




Figure 1: Architecture Dashboard Visualizing HMM Macro Regimes, Total Portfolio PnL vs Benchmark, and Rolling Sharpe Stability.

and four S&P 500 Sector ETFs (XLK, XLF, XLV, XLE). Signals are executed Market-on-Close (MOC), exposing the portfolio to overnight gap risk. A transaction cost of 5 basis points (0.05%) was deducted per unit of turnover.

5 Empirical Results

5.1 Convexity and the Capital-Constrained Mandate

Model A is not a return-maximization strategy; it is a convexity-maximization strategy for tail-risk hedging. Its absolute return (40.27%) underperforms SPY (73.81%), but its Calmar Ratio of 21.65 versus SPY's 2.90 demonstrates its utility for institutional capital preservation. During the 2022 secular decline, market signals lacked institutional backing; the engine safely held the portfolio in cash as shown in Figure 1.

```
=====
V10 OPTIMAL METRICS (NET OF 5 BPS FRICTION)
=====
[*] Net Portfolio PnL:      40.27%
[*] Maximum Drawdown:     1.86%
[*] Sharpe Ratio:          0.22
[*] Sortino Ratio:         5.11
[*] Return Skewness:      +35.34
[*] Veto Rate:             99.8%
```

Figure 2: Console output demonstrating the raw execution metrics of Model A, achieving a Sortino Ratio of 5.11.

5.2 Statistical Robustness and Extreme Skewness

The raw terminal metrics (Figure 2) highlight an extreme positive Return Skewness (+35.34). Over the 5-year sample, Model A executed exactly 12 trades. Consequently, the daily return distribution is heavily concentrated at 0.00%, accompanied by a sparse, fat right tail representing the 12 verified institutional


breakouts.

5.3 Benchmarking VPIN vs. RVWS (Ablation Study)

While VPIN reduced maximum drawdown to 14.20%, RVWS demonstrated superior performance. VPIN frequently misclassified retail-driven “meme-stock” anomalies as informed flow due to volumetric imbalance. RVWS, by calculating *capital cost* of spread absorption, successfully vetoed these toxic events.

5.4 Threshold (τ) Sensitivity Analysis

To verify robustness, a systematic parameter sweep was conducted (Figure 3 and Table 4).



SENSITIVITY TABLE RESULTS			
Tau: 0.30	PnL: -28.95%	Max DD: 38.57%	Veto Rate: 83.5%
Tau: 0.45	PnL: 36.89%	Max DD: 9.30%	Veto Rate: 99.0%
Tau: 0.55	PnL: 40.27%	Max DD: 1.86%	Veto Rate: 99.8%
Tau: 0.60	PnL: 40.27%	Max DD: 1.86%	Veto Rate: 99.8%
Tau: 0.70	PnL: 0.00%	Max DD: 0.00%	Veto Rate: 100.0%

Figure 3: Threshold Sensitivity Sweep demonstrating the phase transition of convexity.

Table 4: Threshold (τ) Impact on Convexity

Threshold (τ)	Net Return	Max DD	Veto Rate
0.30	-28.95%	38.57%	83.5%
0.45	36.89%	9.30%	99.0%
0.55	40.27%	1.86%	99.8%
0.60 (Strict)	40.27%	1.86%	99.8%
0.70	0.00%	0.00%	100.0%

5.5 Transaction Cost Sensitivity

The architecture is highly insulated against execution friction due to low turnover, maintaining convexity even under 10 bps slippage (Table 5).

Table 5: Execution Friction Sensitivity

Friction Tier	Net Return	Max DD	Sortino
0 bps (Theoretical)	40.81%	1.84%	5.14
1 bps (Prime Broker)	40.70%	1.85%	5.13
5 bps (Institutional)	40.27%	1.86%	5.11
10 bps (High Slippage)	39.73%	1.88%	5.05

6 Limitations and Further Research

The institutional order-flow proxy is a heuristic derivation lacking ground-truth labels found in proprietary Level-3 datasets. Furthermore, the RVWS metric assumes moving the price requires prohibitive capital, a property that may violate in micro-cap equities where retail flow can successfully mimic institutional cost signatures. Future research will test this property using Level-2 limit order book tick data.

7 Conclusion

This paper demonstrates that signal generation is secondary to rigorous validation. By layering a Spence-Harsanyi Game Theory veto module beneath an HMM router, we created an architecture capable of mitigating exposure to retail toxicity. This modular framework can be applied as an overlay filter to existing institutional strategies to systematically quarantine execution during periods of high retail toxicity.

References

- [1] Ang, A., & Bekaert, G. (2002). International asset allocation with regime shifts. *The Review of Financial Studies*, 15(4), 1137–1187.
- [2] Easley, D., López de Prado, M. M., & O’Hara, M. (2011). The microstructure of the “Flash Crash”. *The Journal of Portfolio Management*, 37(2), 118–128.
- [3] Efron, B. (1979). Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1), 1–26.
- [4] Glosten, L. R., & Milgrom, P. R. (1985). Bid, ask and transaction prices in a specialist market. *Journal of Financial Economics*, 14(1), 71–100.
- [5] Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series. *Econometrica*, 57(2), 357–384.
- [6] Harsanyi, J. C. (1967). Games with incomplete information played by “Bayesian” players. *Management Science*, 14(3), 159–182.
- [7] Kim, Y., Lee, J., & Lee, J. (2019). Asset allocation with hidden Markov models. *Journal of Risk and Financial Management*, 12(3), 131.

- [8] Kyle, A. S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315–1335.
- [9] Spence, M. (1973). Job market signaling. *The Quarterly Journal of Economics*, 87(3), 355–374.